

Relational Algebra

Miao Qiao

The University of Auckland



Relational Algebra

- Simplest query: **relation name**

instructor

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

(a) The *instructor* table

- Use **operators** to filter, slice, combine

Select operator: pick certain rows

- Select *instructors* where the instructor is in the “Physics” department.

$$\sigma_{dept_name = \text{“Physics”}}(instructor)$$

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 33456 | Gold | Physics | 87000 |

- Select instructors in Physics depart
with salary > 90000.

$$\sigma_{dept_name = \text{“Physics”} \wedge salary > 90000}(instructor)$$

- $\sigma_{cond} Rel$
- Comparisons: =, ≠, >, ≥, <, ≤
- Logic connectives:
 \wedge (and), \vee (or), \neg (not)

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

(a) The *instructor* table

Project operator: pick certain columns

- Pick ID, name and salary of *instructor*

$$\Pi_{ID, name, salary} (instructor)$$

- Result:

| <i>ID</i> | <i>name</i> | <i>salary</i> |
|-----------|-------------|---------------|
| 10101 | Srinivasan | 65000 |
| 12121 | Wu | 90000 |
| 15151 | Mozart | 40000 |
| 22222 | Einstein | 95000 |
| 32343 | El Said | 60000 |
| 33456 | Gold | 87000 |
| 45565 | Katz | 75000 |
| 58583 | Califieri | 62000 |
| 76543 | Singh | 80000 |
| 76766 | Crick | 72000 |
| 83821 | Brandt | 92000 |
| 98345 | Kim | 80000 |

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

$$\Pi_{A_1, A_2, A_3 \dots A_k} (r)$$

Compose operator: pick both rows and columns

- Find the names of all instructors in the Physics department.

$$\sigma_{dept_name = \text{"Physics"}} (instructor)$$

$$\Pi_{name}(\sigma_{dept_name = \text{"Physics"}} (instructor))$$

- $\Pi_{A1, A2} Expr$
- $\sigma_{cond} Expr$

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

(a) The *instructor* table

Cross-product: combines two relations (a.k.a Cartesian-product)

instructor X *teaches*

instructor.ID

| <i>instructor</i> |
|-------------------|
| <u><i>ID</i></u> |
| <i>name</i> |
| <i>dept_name</i> |
| <i>salary</i> |

| <i>teaches</i> |
|-------------------------|
| <u><i>ID</i></u> |
| <u><i>course_id</i></u> |
| <u><i>sec_id</i></u> |
| <u><i>semester</i></u> |
| <u><i>year</i></u> |

teaches.id

| <i>instructor.ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> | <i>teaches.ID</i> | <i>course_id</i> | <i>sec_id</i> | <i>semester</i> | <i>year</i> |
|----------------------|-------------|------------------|---------------|-------------------|------------------|---------------|-----------------|-------------|
| 10101 | Srinivasan | Comp. Sci. | 65000 | 10101 | CS-101 | 1 | Fall | 2017 |
| 10101 | Srinivasan | Comp. Sci. | 65000 | 10101 | CS-315 | 1 | Spring | 2018 |
| 10101 | Srinivasan | Comp. Sci. | 65000 | 10101 | CS-347 | 1 | Fall | 2017 |
| 10101 | Srinivasan | Comp. Sci. | 65000 | 12121 | FIN-201 | 1 | Spring | 2018 |
| 10101 | Srinivasan | Comp. Sci. | 65000 | 15151 | MU-199 | 1 | Spring | 2018 |
| 10101 | Srinivasan | Comp. Sci. | 65000 | 22222 | PHY-101 | 1 | Fall | 2017 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 12121 | Wu | Finance | 90000 | 10101 | CS-101 | 1 | Fall | 2017 |
| 12121 | Wu | Finance | 90000 | 10101 | CS-315 | 1 | Spring | 2018 |
| 12121 | Wu | Finance | 90000 | 10101 | CS-347 | 1 | Fall | 2017 |
| 12121 | Wu | Finance | 90000 | 12121 | FIN-201 | 1 | Spring | 2018 |

Cross-product: combine two relations (a.k.a Cartesian-product)

instructor X *teaches*

- Find the instructors and the courses that they taught

| <i>instructor</i> | <i>teaches</i> |
|-------------------|------------------|
| <u>ID</u> | <u>ID</u> |
| name | <u>course_id</u> |
| dept_name | <u>sec_id</u> |
| salary | <u>semester</u> |
| | <u>year</u> |

$$\sigma_{instructor.id = teaches.id}(instructor \times teaches)$$

Natural join

- Combine two relations, enforce equality on all attributes with same name
- Eliminate the copy of duplicated attributes

| <i>instructor</i> |
|-------------------|
| <u>ID</u> |
| name |
| dept_name |
| salary |

instructor ⋈ *teaches*

| <i>teaches</i> |
|------------------|
| <u>ID</u> |
| <u>course_id</u> |
| <u>sec_id</u> |
| <u>semester</u> |
| <u>year</u> |

- Find all the names of instructors whose department building is 303 and who have taught a course in 2024.

| <i>department</i> |
|-------------------|
| <u>dept_name</u> |
| building |
| budget |

$\Pi_{name}(\sigma_{building = '303' \wedge year = 2024} instructor \bowtie teaches \bowtie department)$

Theta join

- $\text{Exp}_1 \bowtie_{\theta} \text{Exp}_2$
 - θ denotes the selection condition
 - Equivalent to $\sigma_{\theta}(\text{Exp}_1 \bowtie \text{Exp}_2)$
- Basic operation implemented in DBMS

Wrap up

- Data definition language
- Basic steps in creating and using relational DB
- Data manipulation language
- Relational algebra
 - Simplest query: relation name
 - Use operators to filter, slice, combine
 - Operators so far: select, project, cross-product, natural join, theta join

Query Quest

employee (ID, person_name, street, city)

works (ID, person_name, company_name, salary)

company (company_name, city)

- Find the Query
 - a. Find the ID and name of each employee who works for “BigBank”.
 - b. Find the ID, name, and city of residence of each employee who works for “BigBank”.
 - c. Find the ID, name, street address, and city of residence of each employee who works for “BigBank” and earns more than \$10000 per year.

Set operation

■ Union Operation

- Find the names of instructors and the names of the departments

$$\Pi_{name} instructor \cup \Pi_{dept_name} instructor$$

■ Intersection operation

- Find names that are both an instructor name and a department name

$$\Pi_{name} instructor \cap \Pi_{dept_name} department$$

| <i>instructor</i> | <i>department</i> |
|-------------------|-------------------|
| <u>ID</u> | <u>dept_name</u> |
| name | building |
| dept_name | budget |
| salary | |

| ID | name | dept_name | salary |
|-------|------------|------------|--------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

■ Set difference operation

- Find the names of the departments

who has no instructors

$$\Pi_{dept_name} department - \Pi_{dept_name} instructor$$

The Rename Operation

- Rename operator, ρ , name and format the results of an expression
- $\rho_{x(A1,A2, \dots An)}(E)$ general form
 - $\rho_x(E)$
 - $\rho_{A1,A2, \dots An}(E)$
- Functions:
 - To unify schemas for set operations
 - Find names that are both an instructor name and a department name

$$\Pi_{name} \text{instructor} \cap \Pi_{dept_name} \text{instructor}$$

$$\rho_{x(A)}(\Pi_{name} \text{instructor}) \cap \rho_{x(A)}(\Pi_{dept_name} \text{instructor})$$

The Rename Operation

- Rename operator, ρ , name and format the results of an expression

- General form $\rho_{x(A1,A2, \dots An)}(E)$

- $\rho_x(E)$
- $\rho_{A1,A2, \dots An}(E)$

- Functions:

- Unify schemas for set operations
- For disambiguation in “self-joins”
 - Find pairs of instructors in the same dept
 - Instructor x instructor ?

| ID | name | dept_name | salary |
|-------|------------|------------|--------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

(a) The *instructor* table

- $\rho_{r1(a1,b1,c1,d1)} \text{ Instructor } \bowtie \rho_{r2(a2,b2,c2,d2)} \text{ Instructor}$
- $\sigma_{r1.c1 = r2.c2}$
- $\rho_{r1(a1,b1,c,d1)} \text{ Instructor } \bowtie \rho_{r2(a2,b2,c,d2)} \text{ Instructor}$

Assignment Operation

- Break down relational algebra expressions to their parts
- Find all instructor in the “Physics” and Music department.

$Physics \leftarrow \sigma_{dept_name = "Physics"}(instructor)$

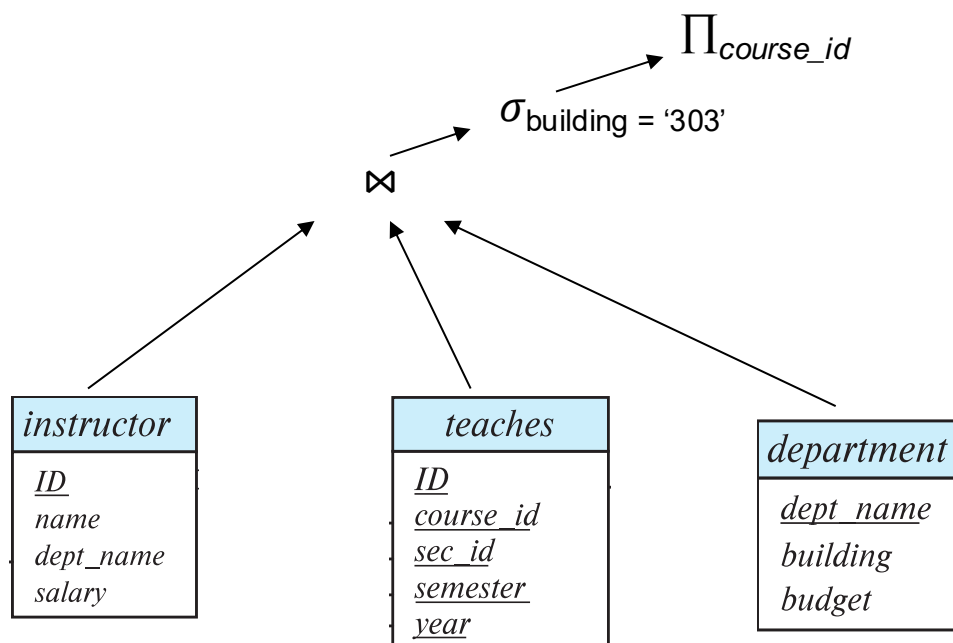
$Music \leftarrow \sigma_{dept_name = "Music"}(instructor)$

$Physics \cup Music$

- The assignment operation is denoted by \leftarrow and works like assignment in a programming language.

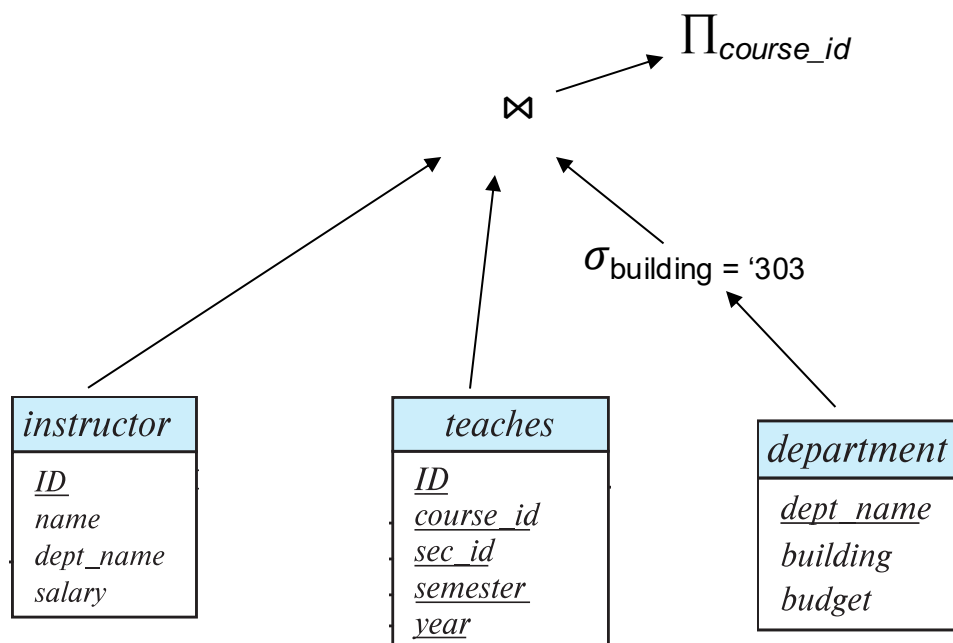
Expression tree

- Find the course IDs whose instructors are in building 303



Expression tree

- Find the course IDs whose instructors are in building 303



- The two queries are not identical; they are, however, equivalent -- they give the same result on any database.

Relational Algebra

- A query language of a set of operations that take one or two relations as input and produce a new relation as their result.
- Operators
 - select: σ
 - project: Π
 - union: \cup
 - set difference: $-$
 - Cartesian product: \times
 - Natural Join \bowtie
 - Theta join \bowtie_{θ}
 - rename: ρ
 - assignment : \leftarrow
- Expression tree

Wrap up

- Relational algebra
 - Simplest query: relation name
 - Use operators to filter, slice, combine
 - Operators: select, project, cross-product, natural join, theta join
 - Set operation, assignment, rename
 - Expression tree

Query Quest

employee (ID, person_name, street, city)

works (ID, person_name, company_name, salary)

company (company_name, city)

- Find the Query
 - Find the ID and name of each employee in this database who lives in the same city as the company for which she or he works.

FIN

Any questions?