

Assignment 2

Due Date: April 18, 2025, at 11:59 pm
10% of the final grade

NAME: _____

UPI: _____

ID: _____

Please ensure that you fill in your name, UPI, and ID above.

It's best to answer Q2 after completing your Week 5 lab, and Q3 after completing your Week 6 lab.

Completely fill the circles as shown: ○○●○

Q1 [3 marks] Assume that each block is 1KB, and the buffer pool is 1MB.

1. [1 mark] What is the maximum file size that can be sorted using external memory sorting in two passes? Recall that the first pass follows the create-runs step.
 - ☒ a. 1TB
 - ☐ b. 10GB
 - ☐ c. 1GB
 - ☐ d. 100MB
 - ☐ e. None of the above
2. [1 mark] Consider a 10GB relation and construct a B+ tree for the relation using external memory sorting. Assume that all leaf nodes (each stored in a block) are full and that all internal nodes reside in the main memory. What is the total number of I/Os required to construct the B+ tree? Note:
 - i. The leaf nodes of the B+ tree are **stored** on disk.
 - ii. Constructing the first-level index of the B+ tree requires reading the sorted file of the relation once.
 - ☒ a. 70M
 - ☐ b. 80M
 - ☐ c. 90M
 - ☐ d. 100M
 - ☐ e. None of the above
3. [1 mark] A colleague is preparing to interview a candidate and would like to assess their knowledge of indexing. The colleague has compiled a list of statements and requests that you identify which ones are true. Please select all the correct statements.
 - ☒ Write-optimized indices can significantly reduce the cost of inserts, and to a lesser extent, of updates, as compared to B+trees. On the other hand, the index lookup cost can be significantly higher for write-optimized indices as compared to B+trees.
 - ☐ NULL values can be easily treated because they represent the absence of a value, making it straightforward to handle them without requiring special attention.
 - ☒ Bloom filters can eliminate unnecessary disk I/Os.
 - ☐ Bloom filters are effective for exact-match (or lookup) queries.
 - ☐ None of the above

Marking Rubric: Selecting options 1, 3, and 4 also earns full marks.

Q2 [4 marks] Consider the B^+ -tree shown in Figure 1 with two levels of nodes. Each leaf node occupies a block, as does each internal node. Each block can hold up to 3 tuples, or alternatively, it can function as an internal node with a fanout of 5. Answer the following questions.

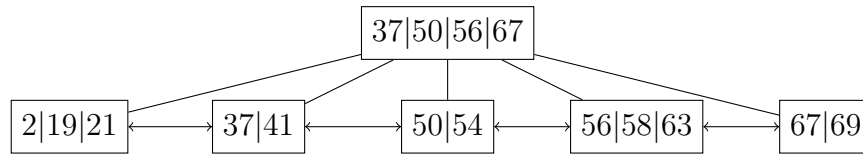


Figure 1: B^+ -tree

1. [1 mark] Insert 59 to the B^+ tree of Figure 1. Select the resulting tree.

- ☐ a.
- ☐ b.
- ☒ c.
- ☐ d.
- ☐ e. None of the above

2. [1 mark] After inserting 59 to Figure 1, how many I/Os are need to find all the tuples with keys in an open range of (35, 57)? Assume the root is stored in the main memory and exclude the I/O cost of reporting the output.

- ☐ a. 2
- ☐ b. 3
- ☐ c. 4
- ☒ d. 5
- ☐ e. None of the above

3. [1 mark] During the sequential insertion of the values 59, 4, 8, 1, 5, and 7 into the B+ Tree shown in Figure 1, how many nodes were split?

- ☐ a. 2
- ☐ b. 3
- ☐ c. 4
- ☐ d. 5
- ☒ e. None of the above

4. [1 mark] How many tuples a B^+ tree of 4 levels can store at most under current parameters?
- ☐ a. 75
 - ☒ b. 375
 - ☐ c. 1875
 - ☐ d. 150
 - ☐ e. None of the above

Q3 [3 marks] Consider an extendible hashing structure such that:

- Each bucket can hold up to three records.
- The hashing function uses the highest g bits (left bits are high) of the hashing value, where g is the global depth.
- A new extendible hashing structure is initialized with $g = 0$ and one empty bucket.
- If multiple keys are provided in a question, assume they are inserted one after the other from left to right.
- Records with duplicate keys will be retained without deduplication.

Key	Hashing value	Key	Hashing value
7	000	18	100
20	110	25	100
37	010	44	010
49	000	50	001
51	010	69	110

1. [1 mark] After inserting 50, 44, 25, 20, 37, what are the local depth of the bucket containing 25 and the global depth, respectively?
- ☐ a. 25 (depth 0). The global depth is 1.
 - ☒ b. 25 (depth 1). The global depth is 1.
 - ☐ c. 25 (depth 2). The global depth is 1.
 - ☐ d. 25 (depth 1). The global depth is 2.
 - ☐ e. None of the above
2. [1 mark] Starting from the results of question Q3.1, insert 7. What are the local depths of the buckets for each key?
- ☒ a. 7 (depth 2), 20 (depth 1), 25 (depth 1), 37 (depth 2).
 - ☐ b. 7 (depth 2), 20 (depth 2), 25 (depth 1), 37 (depth 2).
 - ☐ c. 7 (depth 2), 20 (depth 1), 25 (depth 2), 37 (depth 2).
 - ☐ d. 7 (depth 1), 20 (depth 1), 25 (depth 1), 37 (depth 2).
 - ☐ e. None of the above
3. [1 mark] Starting from the results of question Q3.2, insert 51, 49, 18, 69, 37. What are the local depths of the buckets for each key? What is the global depth?
- ☐ a. 20 (depth 2), 25 (depth 2), 37 (depth 3), 50 (depth 2). Global depth: 3.
 - ☐ b. 20 (depth 2), 25 (depth 2), 37 (depth 3), 50 (depth 3). Global depth: 3.
 - ☒ c. 20 (depth 2), 25 (depth 2), 37 (depth 2), 50 (depth 2). Global depth: 2.
 - ☐ d. 20 (depth 2), 25 (depth 1), 37 (depth 3), 50 (depth 2). Global depth: 3.
 - ☐ e. None of the above