

Introduction to SQL

Miao Qiao

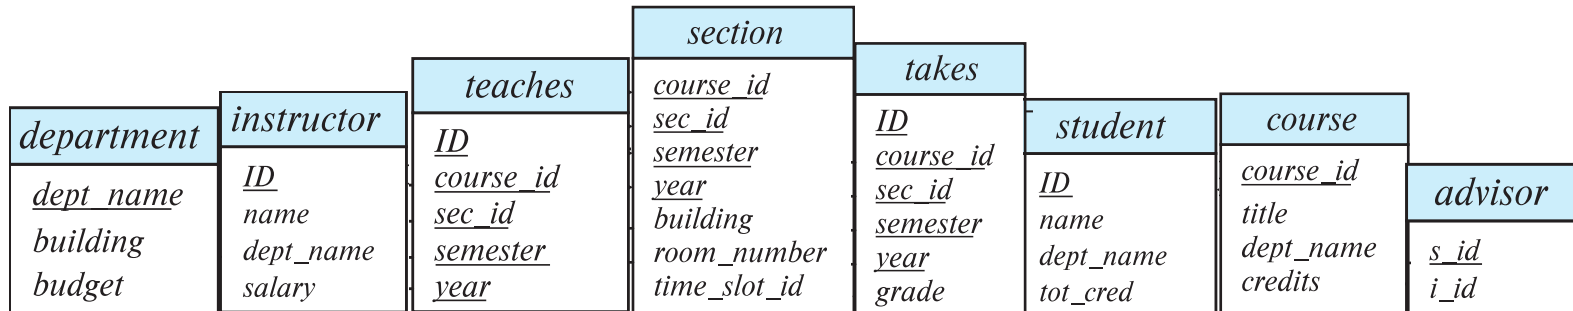
The University of Auckland



Basic Query Structure

- A typical SQL query has the form:

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P



Basic Query Structure (demo)

1. Find the names of all instructors
2. Find the department names of all instructors, and remove duplicates
3. Find the department names of all instructors, not removing duplicates
4. Find all attributes of instructor show the entire instructor table
5. Find a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12
6. Find all instructors in Comp. Sci. dept with salary > 70000
7. Find the names of all instructors who have taught some course and the course_id
8. Find the names of all instructors in the Comp. Sci. department who have taught some course and the course_id
9. Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

department			section		takes		student		course		advisor	
instructor			teaches		ID		student		course		advisor	
ID			course_id		course_id		ID		course_id		s_id	
name			sec_id		sec_id		name		title		i_id	
dept_name			semester		semester		dept_name		dept_name			
building			year		year		tot_cred		credits			
budget			building		grade							
			room_number									
			time_slot_id									

Basic Query Structure (demo)

- Find the names of all instructors whose name includes the substring “in”.
- String Operations
 - The operator **like** uses patterns that are described using two special characters:
 - percent (%). The % character matches any substring.
 - underscore (_). The _ character matches any character
- Find the names of all instructors whose name has 4 characters.
- Find the names of all instructors whose name has at least 4 characters.

				section				
department	instructor	teaches			takes	student	course	
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	
building	name	<u>sec_id</u>	<u>semester</u>	<u>sec_id</u>	<u>course_id</u>	name	title	advisor
budget	dept_name	<u>year</u>	<u>building</u>	<u>year</u>	<u>sec_id</u>	dept_name	dept_name	<u>s_id</u>
	salary		<u>room_number</u>	<u>time_slot_id</u>	<u>semester</u>	tot_cred	credits	<u>i_id</u>
					<u>year</u>			
					grade			

Basic Query Structure (demo)

1. List in alphabetic order the names of all instructors
2. List in descending alphabetic order the names of all instructors
3. List in order of the combination of the names and salary of all instructors
4. Find the names of all instructors with salary between \$90,000 and \$100,000
5. Find courses that ran in Fall 2017 or in Spring 2018
6. Find courses that ran in Fall 2017 and in Spring 2018
7. Find courses that ran in Fall 2017 but not in Spring 2018
8. Find courses that ran in Fall 2017 or in Spring 2018, retain all duplications
9. Find all instructors whose salary is null
10. Find all instructors whose salary is not null
11. Null under and, or, with true/false

		<i>teaches</i>		<i>section</i>	<i>takes</i>		<i>course</i>	<i>advisor</i>
<i>department</i>	<i>instructor</i>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>course_id</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>course_id</i></u>	
<u><i>dept_name</i></u>	<u><i>ID</i></u>	<i>sec_id</i>	<i>semester</i>	<i>sec_id</i>	<i>sec_id</i>	<i>semester</i>	<i>title</i>	<u><i>s_id</i></u>
<i>building</i>	<i>name</i>	<i>year</i>	<i>building</i>	<i>room_number</i>	<i>year</i>	<i>grade</i>	<i>dept_name</i>	<i>i_id</i>
<i>budget</i>	<i>dept_name</i>		<i>time_slot_id</i>				<i>credits</i>	
	<i>salary</i>							

Basic Query Structure (demo)

- These functions operate on the multiset of values of a column of a relation, and return a value

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

select A_1, A_2, \dots, A_n  Aggregation function over values over multiple rows

from r_1, r_2, \dots, r_m

where P

group by columns  New clauses

having condition 

department		instructor		teaches	section	takes	student	course	advisor
<u>dept_name</u>		<u>ID</u>		<u>ID</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	
building		name		<u>course_id</u>	<u>sec_id</u>	<u>course_id</u>	name	title	<u>s_id</u>
		dept_name		<u>sec_id</u>	<u>semester</u>	<u>sec_id</u>	dept_name	dept_name	<u>i_id</u>
budget		salary		<u>semester</u>	building	<u>semester</u>	tot_cred	credits	
				<u>year</u>	room_number	<u>year</u>			
					time_slot_id	grade			

Basic Query Structure (demo)

- These functions operate on the multiset of values of a column of a relation, and return a value

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

- Find the highest salary of any instructor
- Find the average salary of instructors in the Computer Science department
- Find the lowest salary of an instructor who have taught a course
- Find the total number of instructors who teach a course in the Spring 2018 semester
- Find the number of tuples in the *course* relation

		<i>teaches</i>		<i>takes</i>			
<i>department</i>	<i>instructor</i>	<i>section</i>		<i>student</i>		<i>course</i>	<i>advisor</i>
<u>dept_name</u> building budget	<u>ID</u> name dept_name salary	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>	<u>course_id</u> <u>sec_id</u> <u>semester</u> year building room_number time_slot_id	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> grade	<u>ID</u> name dept_name tot_cred	<u>course_id</u> title dept_name credits	<u>s_id</u> <u>i_id</u>

Aggregate – Group By - Having

1. Find the average salary of instructors in each department
2. Find the names and average salaries of all departments whose average salary is greater than 42000
3. Find the names and average salaries of all departments over instructors whose salary is greater than 70000
4. Find the names and average salaries of all departments whose average salary is greater than 70000
5. Find the average salaries of instructors who have taught a course

				section				
department	instructor	teaches			takes	student	course	
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>advisor</u>
building	name	<u>sec_id</u>	<u>sec_id</u>	<u>sec_id</u>	<u>course_id</u>	name	title	<u>s_id</u>
	dept_name	<u>semester</u>	<u>semester</u>	<u>semester</u>	<u>sec_id</u>	dept_name	dept_name	<u>i_id</u>
budget	salary	<u>year</u>	<u>year</u>	<u>year</u>	<u>semester</u>	tot_cred	credits	
				<u>building</u>	<u>year</u>			
				<u>room_number</u>	<u>grade</u>			
				<u>time_slot_id</u>				

Nested Subqueries

- A **subquery** is a **select-from-where** expression that is nested within another query.
- The nesting can be done in the following SQL query

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

as follows:

- **From clause:** r_i can be replaced by any valid subquery
- **Where clause:** P can be replaced with an expression of the form:

$B <\text{operation}> (\text{subquery})$

B is an attribute and $<\text{operation}>$ to be defined later.

- **Select clause:**

A_i can be replaced by a subquery that generates a single value.

		teaches		section	takes		course	advisor
department	instructor	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>	<u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> <u>building</u> <u>room_number</u> <u>time_slot_id</u>		<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> <u>grade</u>	<u>ID</u> <u>name</u> <u>dept_name</u> <u>tot_cred</u>	<u>course_id</u> <u>title</u> <u>dept_name</u> <u>credits</u>	
<u>dept_name</u> building budget	<u>ID</u> name dept_name salary							<u>s_id</u> <u>i_id</u>

Subquery in Where Clause

1. Name all instructors whose name is either “Mozart” or Einstein”
2. Name all instructors whose name is neither “Mozart” nor Einstein”
3. Find the total number of (distinct) students who have taken course sections taught by the instructor with *ID* 10101 or 12121 or 15151
4. Find names of instructors with salary greater than that of some (at least one) instructor in the Computer Science department.
5. Find all instructors earning the highest salary (there may be more than one with the same salary).
6. Find the names of all instructors whose salary is greater than the salary of all instructors in the Biology department.

		<i>teaches</i>		<i>section</i>	<i>takes</i>		<i>student</i>	<i>course</i>	<i>advisor</i>
<i>department</i>	<i>instructor</i>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>sec_id</i></u>	<u><i>course_id</i></u>	<u><i>sec_id</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	
<u><i>dept_name</i></u>	<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>course_id</i>	<i>sec_id</i>	<i>name</i>	<i>title</i>	<u><i>s_id</i></u>
<i>building</i>	<i>dept_name</i>	<i>semester</i>	<i>building</i>	<i>room_number</i>	<i>semester</i>	<i>year</i>	<i>dept_name</i>	<i>dept_name</i>	<i>i_id</i>
<i>budget</i>	<i>salary</i>	<i>year</i>	<i>time_slot_id</i>	<i>grade</i>	<i>tot_cred</i>			<i>credits</i>	

Definition of “some” Clause

- $F <\text{comp}> \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F <\text{comp}> t)$

Where $<\text{comp}>$ can be: $<$, \leq , $>$, $=$, \neq

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$ (read: 5 < some tuple in the relation)

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$

However, $(\neq \text{some}) \not\equiv \text{not in}$

Definition of “all” Clause

- $F \text{ <comp> all } r \Leftrightarrow \forall t \in r (F \text{ <comp> } t)$

$$(5 < \text{all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$$

$$(5 < \text{all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$$

$$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$$

$(\neq \text{all}) \equiv \text{not in}$

However, $(= \text{all}) \not\equiv \text{in}$

Subquery in Where Clause

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
 - Find all courses taught in both the Fall 2017 semester and in the Spring 2018 semester
- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- The **unique** construct evaluates to “true” if a given subquery contains no duplicates .
 - Find all courses that were offered at most once in 2017

		<i>teaches</i>		<i>section</i>	<i>takes</i>		<i>student</i>	<i>course</i>	<i>advisor</i>
<i>department</i>	<i>instructor</i>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>course_id</i></u>	<u><i>sec_id</i></u>	<u><i>semester</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	
<u><i>dept_name</i></u>	<u><i>ID</i></u>	<i>course_id</i>	<u><i>sec_id</i></u>	<i>semester</i>	<i>course_id</i>	<u><i>sec_id</i></u>	<i>name</i>	<i>title</i>	<u><i>s_id</i></u>
<i>building</i>	<i>dept_name</i>	<u><i>semester</i></u>	<i>building</i>	<i>room_number</i>	<u><i>semester</i></u>	<i>year</i>	<i>dept_name</i>	<i>dept_name</i>	<i>i_id</i>
<i>budget</i>	<i>salary</i>	<u><i>year</i></u>	<i>time_slot_id</i>	<i>grade</i>	<i>tot_cred</i>			<i>credits</i>	

Query Quest

1. Find names of instructors with salary greater than that of some (at least one) instructor in the Computer Science department.
 - Use self-join, some, exists, aggregation
2. Find all instructors earning the highest salary (there may be more than one with the same salary).
 - Use self-join, all, exists, aggregation

				section				
department	instructor	teaches			takes	student	course	
<u>dept_name</u> building budget	<u>ID</u> name dept_name salary	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>		<u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> building room_number time_slot_id	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> grade	<u>ID</u> name dept_name tot_cred	<u>course_id</u> title dept_name credits	advisor
								<u>s_id</u> <u>i_id</u>

Subquery in From and Select Clauses

1. Find the average instructors' salaries of those departments where the average salary is greater than \$42,000"
2. Find all departments with the maximum budget (with clause)
3. Find all departments where the total salary is greater than the average of the total salary at all departments
4. List all departments along with the number of instructors in each department

				section				
department	instructor	teaches			takes	student	course	
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>advisor</u>
building	name	<u>sec_id</u>	<u>sec_id</u>	<u>sec_id</u>	<u>course_id</u>	name	title	<u>s_id</u>
budget	dept_name	<u>semester</u>	<u>semester</u>	building	<u>sec_id</u>	dept_name	dept_name	<u>i_id</u>
	salary	<u>year</u>	<u>year</u>	room_number	<u>semester</u>	tot_cred	credits	
				time_slot_id	<u>year</u>			
					grade			

With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

```

with max_budget (value) as
    (select max(budget)
     from department)
select department.name
from department, max_budget
where department.budget = max_budget.value;
  
```

<i>department</i>	<i>instructor</i>	<i>teaches</i>	<i>section</i>	<i>takes</i>	<i>student</i>	<i>course</i>	<i>advisor</i>
<u><i>dept_name</i></u> <i>building</i> <i>budget</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>salary</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u>	<u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>building</i> <i>room_number</i> <i>time_slot_id</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>grade</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>tot_cred</i>	<u><i>course_id</i></u> <i>title</i> <i>dept_name</i> <i>credits</i>	<u><i>s_id</i></u> <u><i>i_id</i></u>

Modification of the Database

- Deletion of tuples from a given relation.
 - Insertion of new tuples into a given relation
 - Updating of values in some tuples in a given relation
-
- **delete from** *a relation* **where** condition

```
delete from instructor  
where dept name in (select dept name  
                        from department  
                        where building = 'Watson');
```

			<i>section</i>				
<i>department</i>	<i>instructor</i>	<i>teaches</i>		<i>takes</i>			
<u><i>dept_name</i></u>	<u><i>ID</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>ID</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	
<i>building</i>	<i>name</i>	<u><i>course_id</i></u>	<u><i>sec_id</i></u>	<u><i>course_id</i></u>	<i>name</i>	<i>title</i>	<i>advisor</i>
<i>budget</i>	<i>dept_name</i>	<u><i>sec_id</i></u>	<u><i>semester</i></u>	<u><i>sec_id</i></u>	<i>dept_name</i>	<i>dept_name</i>	<u><i>s_id</i></u>
	<i>salary</i>	<u><i>semester</i></u>	<i>building</i>	<u><i>year</i></u>	<i>tot_cred</i>	<i>credits</i>	<u><i>i_id</i></u>
		<u><i>year</i></u>	<i>room_number</i>	<i>grade</i>			
			<i>time_slot_id</i>				

Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor
where salary < (select avg (salary)
                from instructor);
```

- Problem: as we delete tuples from *instructor*, the average salary changes
- Solution used in SQL:
 - First, compute **avg** (salary) and find all tuples to delete
 - Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

				section				
department	instructor	teaches			takes	student	course	advisor
<u>dept_name</u> building budget	<u>ID</u> name dept_name salary	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>		<u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> building room_number time_slot_id	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> grade	<u>ID</u> name dept_name tot_cred	<u>course_id</u> title dept_name credits	<u>s_id</u> <u>i_id</u>

Insertion

- Add a new tuple to *course*

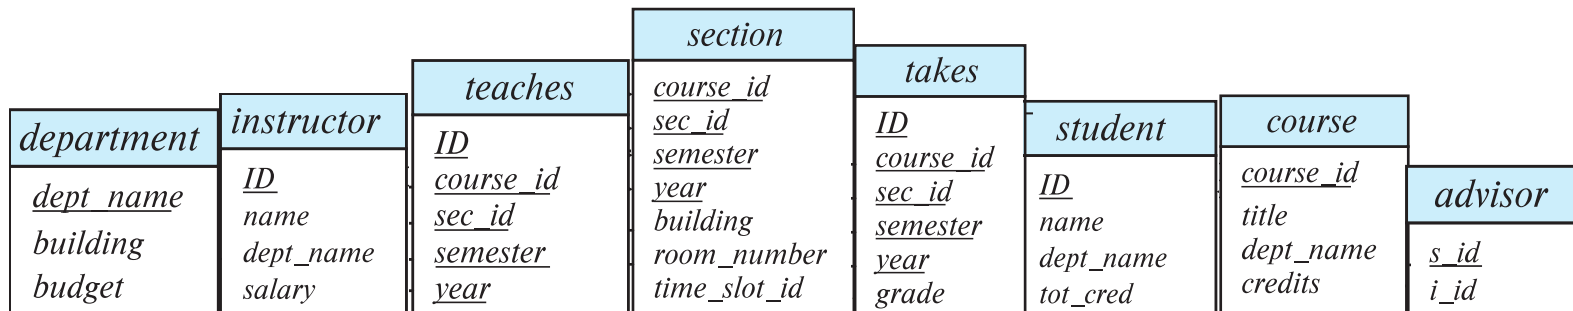
```
insert into course
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently

```
insert into course (course_id, title, dept_name, credits)
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to *student* with *tot_creds* set to null

```
insert into student
values ('3003', 'Green', 'Finance', null);
```



Insertion (Cont.)

- Make each student in the Music department who has earned more than 144 credit hours an instructor in the Music department with a salary of \$18,000.

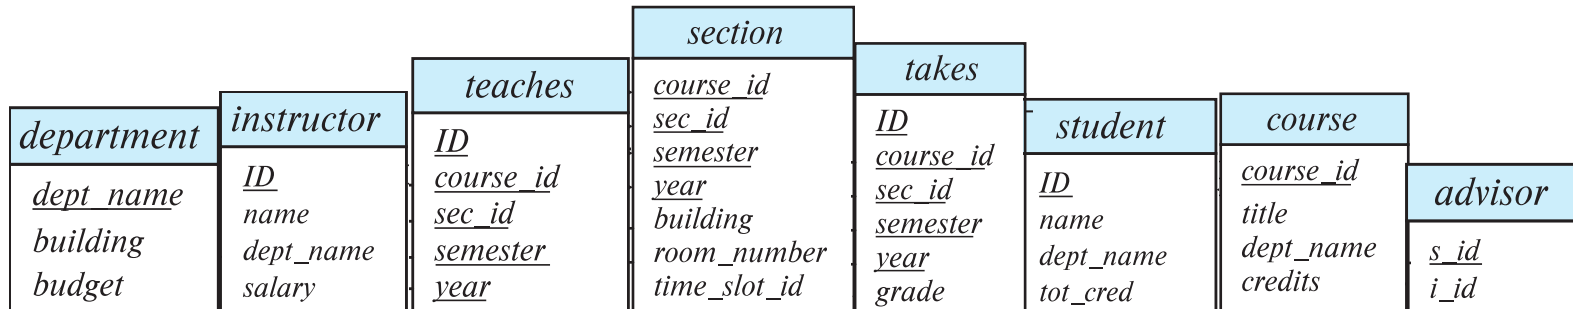
```
insert into instructor
  select ID, name, dept_name, 18000
 from student
 where dept_name = 'Music' and total_cred > 144;
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

Otherwise queries like

```
insert into table1 select * from table1
```

would cause problem



Updates

1. Give a 5% salary raise to all instructors
2. Give a 5% salary raise to those instructors who earn less than 70000
3. Give a 5% salary raise to instructors whose salary is less than average

				section				
department	instructor	teaches			takes	student	course	advisor
<u>dept_name</u> building budget	<u>ID</u> name dept_name salary	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>		<u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> building room_number time_slot_id	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> grade	<u>ID</u> name dept_name tot_cred	<u>course_id</u> title dept_name credits	<u>s_id</u> <u>i_id</u>

Updates (Cont.)

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
 - **update** *instructor*

```

          set salary = salary * 1.03
        where salary > 100000;
        update instructor
          set salary = salary * 1.05
        where salary <= 100000;
```
 - The order is important
 - Can be done better using the **case** statement

update *instructor*

set salary = **case**

when salary <= 100000 **then** salary * 1.05

else salary * 1.03

end

department		instructor		teaches	section	takes	student	course	advisor
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>name</u>	<u>ID</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	
building	dept_name	name	dept_name	course_id	sec_id	course_id	sec_id	title	s_id
budget	salary	dept_name	salary	sec_id	semester	sec_id	semester	dept_name	i_id
		year		year	building	year	year	credits	
					room_number	grade			
					time_slot_id				

Updates with Scalar Subqueries

- Recompute and update `tot_cred` value for all students

update *student* *S*

set *tot_cred* = (**select** **sum**(*credits*)

from *takes*, *course*

where *takes.course_id* = *course.course_id* **and**

S.ID = *takes.ID* **and**

takes.grade <> 'F' **and**

takes.grade **is not null**);

- Sets *tot_cred* to null for students who have not taken any course
 - Instead of **sum**(*credits*), use:

case

when **sum**(*credits*) **is not null** **then** **sum**(*credits*)

else 0

end

		<i>teaches</i>		<i>section</i>	<i>takes</i>		<i>student</i>	<i>course</i>	<i>advisor</i>
<i>department</i>	<i>instructor</i>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>sec_id</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	
<i>dept_name</i>	<i>ID</i>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>course_id</i>	<i>sec_id</i>	<i>name</i>	<i>title</i>	<i>s_id</i>
<i>building</i>	<i>dept_name</i>	<i>semester</i>	<i>year</i>	<i>building</i>	<i>sec_id</i>	<i>semester</i>	<i>dept_name</i>	<i>dept_name</i>	<i>i_id</i>
<i>budget</i>	<i>salary</i>	<i>year</i>	<i>time_slot_id</i>	<i>room_number</i>	<i>year</i>	<i>grade</i>	<i>tot_cred</i>	<i>credits</i>	

Join conditions

- List the names of students instructors along with the titles of courses that they have taken
 - Natural Join with **Using** Clause
 - Join condition with **on** condition

<i>department</i>	<i>instructor</i>	<i>teaches</i>	<i>section</i>	<i>takes</i>	<i>student</i>	<i>course</i>	<i>advisor</i>
<u><i>dept_name</i></u> <i>building</i> <i>budget</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>salary</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u>	<u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>building</i> <i>room_number</i> <i>time_slot_id</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>grade</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>tot_cred</i>	<u><i>course_id</i></u> <i>title</i> <i>dept_name</i> <i>credits</i>	<u><i>s_id</i></u> <u><i>i_id</i></u>

Join types

- Three forms of outer join:
 - Natural left outer join $course \bowtie_{\text{prereq}}$
 - Natural right outer join $course \bowtie_{\text{prereq}}$
 - Natural full outer join $course \bowtie_{\text{prereq}}$

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

<i>department</i>		<i>instructor</i>	<i>teaches</i>	<i>section</i>	<i>takes</i>	<i>student</i>	<i>course</i>	<i>advisor</i>
<u><i>dept_name</i></u>	<i>building</i>	<u><i>ID</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>sec_id</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<i>s_id</i>
<i>budget</i>		<i>name</i>	<i>course_id</i>	<i>semester</i>	<i>course_id</i>	<i>name</i>	<i>title</i>	<i>i_id</i>
		<i>dept_name</i>	<i>sec_id</i>	<i>year</i>	<i>sec_id</i>	<i>dept_name</i>	<i>dept_name</i>	
		<i>salary</i>	<i>semester</i>	<i>building</i>	<i>semester</i>	<i>tot_cred</i>	<i>credits</i>	
			<i>year</i>	<i>room_number</i>	<i>year</i>			
				<i>time_slot_id</i>	<i>grade</i>			

Join

Join types
inner join
left outer join
right outer join
full outer join

Join conditions
natural
on <predicate>
using (A_1, A_2, \dots, A_n)

<u>course_id</u>	<u>title</u>	<u>dept_name</u>	<u>credits</u>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

<u>course_id</u>	<u>prereq_id</u>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

<u>department</u>	<u>instructor</u>	<u>teaches</u>	<u>section</u>	<u>takes</u>	<u>student</u>	<u>course</u>	<u>advisor</u>
<u>dept_name</u> building budget	<u>ID</u> name dept_name salary	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>	<u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> building room_number time_slot_id	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> grade	<u>ID</u> name dept_name tot_cred	<u>course_id</u> title dept_name credits	<u>s_id</u> <u>i_id</u>

Joined Relations – Examples

- course natural right outer join prereq

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

- course full outer join prereq using (*course_id*)

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

<i>department</i>	<i>instructor</i>	<i>teaches</i>	<i>section</i>	<i>takes</i>	<i>student</i>	<i>course</i>	<i>advisor</i>
<u><i>dept_name</i></u> <i>building</i> <i>budget</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>salary</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u>	<u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>building</i> <i>room_number</i> <i>time_slot_id</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>grade</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>tot_cred</i>	<u><i>course_id</i></u> <i>title</i> <i>dept_name</i> <i>credits</i>	<i>s_id</i> <i>i_id</i>

Joined Relations – Examples

- course inner join prereq on
course.course_id = prereq.course_id

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- What is the difference between the above, and a natural join?
- course left outer join prereq on
course.course_id = prereq.course_id

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	null	null

		teaches		takes		course		advisor
department	instructor	ID	course_id	ID	course_id	student	course	
dept_name	ID	course_id	sec_id	course_id	sec_id	ID	course_id	s_id
building	name	sec_id	semester	sec_id	semester	name	title	i_id
budget	dept_name	semester	year	year	year	dept_name	dept_name	
	salary	year	building	grade	tot_cred	credits		
			room_number					
			time_slot_id					

Joined Relations – Examples

- course natural right outer join prereq

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

- course full outer join prereq using (*course_id*)

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

<i>department</i>	<i>instructor</i>	<i>teaches</i>	<i>section</i>	<i>takes</i>	<i>student</i>	<i>course</i>	<i>advisor</i>
<u><i>dept_name</i></u> <i>building</i> <i>budget</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>salary</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u>	<u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>building</i> <i>room_number</i> <i>time_slot_id</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>grade</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>tot_cred</i>	<u><i>course_id</i></u> <i>title</i> <i>dept_name</i> <i>credits</i>	<i>s_id</i> <i>i_id</i>

View

- A **view** provides a mechanism to hide certain data from the view of certain users.
 - Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by


```
select ID, name, dept_name
from instructor
```
 - **create view v as** < query expression >
1. A view of instructors without their salary
 2. Find all instructors in the Biology department
 3. Create a view of department salary totals

				section				
department	instructor	teaches			takes	student	course	advisor
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>s_id</u>
building	name	<u>sec_id</u>	<u>sec_id</u>	<u>semester</u>	<u>course_id</u>	name	title	i_id
budget	dept_name	<u>semester</u>	<u>semester</u>	building	<u>sec_id</u>	dept_name	dept_name	
	salary	<u>year</u>	<u>year</u>	room_number	<u>year</u>	tot_cred	credits	
				time_slot_id	grade			

Define Views using other Views

- create view ***physics_fall_2017*** as
 select *course.course_id*, *sec_id*, *building*, *room_number*
 from *course*, *section*
 where *course.course_id* = *section.course_id*
 and *course.dept_name* = 'Physics'
 and *section.semester* = 'Fall'
 and *section.year* = '2017';
- create view ***physics_fall_2017_watson*** as
 select *course_id*, *room_number*
 from ***physics_fall_2017***
 where *building* = 'Watson';

		<i>teaches</i>		<i>section</i>	<i>takes</i>		<i>course</i>	<i>advisor</i>
<i>department</i>	<i>instructor</i>	<u><i>ID</i></u>		<u><i>course_id</i></u>	<u><i>ID</i></u>	<i>student</i>		
<u><i>dept_name</i></u>	<u><i>ID</i></u>	<i>course_id</i>		<u><i>sec_id</i></u>	<i>course_id</i>	<u><i>ID</i></u>	<u><i>course_id</i></u>	
<i>building</i>	<i>name</i>	<u><i>sec_id</i></u>		<i>semester</i>	<u><i>sec_id</i></u>	<i>name</i>	<i>title</i>	<u><i>s_id</i></u>
<i>budget</i>	<i>dept_name</i>	<u><i>semester</i></u>		<i>building</i>	<u><i>semester</i></u>	<i>dept_name</i>	<i>dept_name</i>	<i>i_id</i>
	<i>salary</i>	<u><i>year</i></u>		<i>room_number</i>	<i>year</i>	<i>tot_cred</i>	<i>credits</i>	
				<i>time_slot_id</i>	<i>grade</i>			

View Expansion

- Expand the view :

```
create view physics_fall_2017_watson as
select course_id, room_number
from physics_fall_2017
where building= 'Watson'
```

- To: create view **physics_fall_2017_watson** as


```
select course_id, room_number
from (select course.course_id, building, room_number
      from course, section
      where course.course_id = section.course_id
        and course.dept_name = 'Physics'
        and section.semester = 'Fall'
        and section.year = '2017')
where building= 'Watson';
```

department		instructor		teaches	section	takes	student	course	advisor
<u>dept_name</u>		<u>ID</u>		<u>ID</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>s_id</u>
building		name		course_id	sec_id	course_id	name	title	i_id
budget		dept_name		sec_id	semester	sec_id	dept_name	dept_name	
		salary		semester	building	semester	tot_cred	credits	
				year	room_number	year			
					time_slot_id	grade			

Views

- Materialized Views
- Update Views
 - Add a new tuple to *faculty* view which we defined earlier
insert into *faculty* values ('30765', 'Green', 'Music');

				section				
department	instructor	teaches			takes	student	course	advisor
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	
building	name	<u>sec_id</u>	<u>sec_id</u>	<u>semester</u>	<u>course_id</u>	name	title	<u>s_id</u>
budget	dept_name	<u>semester</u>	<u>semester</u>	building	<u>sec_id</u>	dept_name	dept_name	<u>i_id</u>
	salary	<u>year</u>	<u>year</u>	room_number	<u>semester</u>	tot_cred	credits	
				time_slot_id	<u>year</u>			
					grade			

Some Updates Cannot be Translated Uniquely

- **create view** *instructor_info* as
 select *ID, name, building*
 from *instructor, department*
 where *instructor.dept_name = department.dept_name;*
- **insert into** *instructor_info* **values** ('69987', 'White', 'Taylor');
- Issues
 - Which department, if multiple departments in Taylor?
 - What if no department is in Taylor?
- Most SQL implementations allow updates only on simple views
 - The **from** clause has only one database relation.
 - The **select** clause contains only attribute names of the relation, and does not have any expressions, aggregates, or **distinct** specification.
 - Any attribute not listed in the **select** clause can be set to null
 - The query does not have a **group** by or **having** clause.

FIN

Any questions?