

Introduction to SQL

Miao Qiao

The University of Auckland



Basic Query Structure

- A typical SQL query has the form:

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m
where P

<i>department</i>	<i>instructor</i>	<i>teaches</i>	<i>section</i>	<i>takes</i>	<i>student</i>	<i>course</i>	<i>advisor</i>
<u><i>dept_name</i></u> <i>building</i> <i>budget</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>salary</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u>	<u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>building</i> <i>room_number</i> <i>time_slot_id</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>grade</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>tot_cred</i>	<u><i>course_id</i></u> <i>title</i> <i>dept_name</i> <i>credits</i>	<u><i>s_id</i></u> <u><i>i_id</i></u>

Basic Query Structure (demo)

1. Find the names of all instructors
2. Find the department names of all instructors, and remove duplicates
3. Find the department names of all instructors, not removing duplicates
4. Find all attributes of instructor show the entire instructor table
5. Find a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12
6. Find all instructors in Comp. Sci. dept with salary > 70000
7. Find the names of all instructors who have taught some course and the course_id
8. Find the names of all instructors in the Comp. Sci. department who have taught some course and the course_id
9. Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

		<i>teaches</i>		<i>section</i>	<i>takes</i>		<i>course</i>	<i>advisor</i>
<i>department</i>	<i>instructor</i>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	
<u>dept_name</u>	<u>ID</u>	<u>course_id</u>	<u>sec_id</u>	<u>sec_id</u>	<u>course_id</u>	<u>sec_id</u>	<u>title</u>	<u>s_id</u>
<i>building</i>	<i>name</i>	<u>semester</u>	<i>building</i>	<i>semester</i>	<i>sec_id</i>	<i>semester</i>	<i>dept_name</i>	<i>i_id</i>
<i>budget</i>	<i>dept_name</i>	<u>year</u>	<i>room_number</i>	<i>year</i>	<i>year</i>	<i>grade</i>	<i>credits</i>	
	<i>salary</i>		<i>time_slot_id</i>					

Basic Query Structure (demo)

- Find the names of all instructors whose name includes the substring “in”.
- String Operations
 - The operator **like** uses patterns that are described using two special characters:
 - percent (%). The % character matches any substring.
 - underscore (_). The _ character matches any character
- Find the names of all instructors whose name has 4 characters.
- Find the names of all instructors whose name has at least 4 characters.

<i>department</i>	<i>instructor</i>	<i>teaches</i>	<i>section</i>	<i>takes</i>	<i>student</i>	<i>course</i>	<i>advisor</i>
<u>dept_name</u> building budget	<u>ID</u> name dept_name salary	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>	<u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> building room_number time_slot_id	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> grade	<u>ID</u> name dept_name tot_cred	<u>course_id</u> title dept_name credits	<u>s_id</u> <u>i_id</u>

Basic Query Structure (demo)

1. List in alphabetic order the names of all instructors
2. List in descending alphabetic order the names of all instructors
3. List in order of the combination of the names and salary of all instructors
4. Find the names of all instructors with salary between \$90,000 and \$100,000
5. Find courses that ran in Fall 2017 or in Spring 2018
6. Find courses that ran in Fall 2017 and in Spring 2018
7. Find courses that ran in Fall 2017 but not in Spring 2018
8. Find courses that ran in Fall 2017 or in Spring 2018, retain all duplications
9. Find all instructors whose salary is null
10. Find all instructors whose salary is not null
11. Null under and, or, with true/false

		<i>teaches</i>		<i>section</i>	<i>takes</i>		<i>course</i>	<i>advisor</i>
<i>department</i>	<i>instructor</i>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>course_id</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>course_id</i></u>	
<u><i>dept_name</i></u>	<u><i>ID</i></u>	<i>course_id</i>	<i>sec_id</i>	<i>semester</i>	<i>course_id</i>	<i>sec_id</i>	<i>title</i>	<u><i>s_id</i></u>
<i>building</i>	<i>name</i>	<i>sec_id</i>	<i>year</i>	<i>building</i>	<i>sec_id</i>	<i>semester</i>	<i>dept_name</i>	<i>i_id</i>
<i>budget</i>	<i>dept_name</i>	<i>semester</i>	<i>room_number</i>	<i>room_number</i>	<i>year</i>	<i>grade</i>	<i>credits</i>	
	<i>salary</i>	<i>year</i>	<i>time_slot_id</i>	<i>time_slot_id</i>				

Basic Query Structure (demo)

- These functions operate on the multiset of values of a column of a relation, and return a value

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

select A_1, A_2, \dots, A_n  Aggregation function over values over multiple rows

from r_1, r_2, \dots, r_m

where P

group by columns  New clauses

having condition 

department		instructor		teaches	section	takes	student	course	advisor
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>name</u>	<u>ID</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	
building	dept_name	name	dept_name	course_id	sec_id	course_id	name	title	s_id
budget	salary	sec_id		semester	year	sec_id	dept_name	dept_name	i_id
		year		building	room_number	semester	tot_cred	credits	
				time_slot_id		year			
						grade			

Basic Query Structure (demo)

- These functions operate on the multiset of values of a column of a relation, and return a value

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

- Find the highest salary of any instructor.
- Find the average salary of instructors in the Computer Science department
- Find the lowest salary of an instructor who have taught a course
- Find the total number of instructors who teach a course in the Spring 2018 semester
- Find the number of tuples in the *course* relation

		<i>teaches</i>		<i>section</i>	<i>takes</i>		<i>student</i>	<i>course</i>	<i>advisor</i>
<i>department</i>	<i>instructor</i>	<u><i>ID</i></u>	<u><i>course_id</i></u>	<u><i>course_id</i></u>	<u><i>sec_id</i></u>	<u><i>semester</i></u>	<u><i>ID</i></u>	<u><i>course_id</i></u>	
<u><i>dept_name</i></u>	<u><i>ID</i></u>	<i>name</i>	<i>sec_id</i>	<i>semester</i>	<i>building</i>	<i>room_number</i>	<i>name</i>	<i>title</i>	<u><i>s_id</i></u>
<i>building</i>	<i>dept_name</i>	<u><i>semester</i></u>	<i>year</i>	<i>time_slot_id</i>	<i>year</i>	<i>grade</i>	<i>dept_name</i>	<i>credits</i>	<u><i>i_id</i></u>
<i>budget</i>	<i>salary</i>						<i>tot_cred</i>		

Aggregate – Group By - Having

1. Find the average salary of instructors in each department
2. Find the names and average salaries of all departments whose average salary is greater than 42000
3. Find the names and average salaries of all departments over instructors whose salary is greater than 7000
4. Find the names and average salaries of all departments whose average salary is greater than 70000

				section				
department	instructor	teaches			takes	student	course	
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>advisor</u>
building	name	<u>sec_id</u>	<u>sec_id</u>	<u>sec_id</u>	<u>course_id</u>	name	title	<u>s_id</u>
	dept_name	<u>semester</u>	<u>semester</u>	<u>year</u>	<u>sec_id</u>	dept_name	dept_name	<u>i_id</u>
budget	salary	<u>year</u>	<u>year</u>	building	<u>semester</u>	tot_cred	credits	
				room_number	<u>year</u>			
				time_slot_id	grade			

Nested Subqueries

- A **subquery** is a **select-from-where** expression that is nested within another query.
- The nesting can be done in the following SQL query

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

as follows:

- **From clause:** r_i can be replaced by any valid subquery
- **Where clause:** P can be replaced with an expression of the form:

$B <\text{operation}> (\text{subquery})$

B is an attribute and $<\text{operation}>$ to be defined later.

- **Select clause:**

A_i can be replaced by a subquery that generates a single value.

		teaches		section	takes		course	advisor
department	instructor	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>	<u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> <u>building</u> <u>room_number</u> <u>time_slot_id</u>		<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> <u>grade</u>	<u>ID</u> <u>name</u> <u>dept_name</u> <u>tot_cred</u>	<u>course_id</u> <u>title</u> <u>dept_name</u> <u>credits</u>	
<u>dept_name</u> building budget	<u>ID</u> name dept_name salary							<u>s_id</u> <u>i_id</u>

<i>department</i>	<i>instructor</i>	<i>teaches</i>	<i>section</i>	<i>takes</i>	<i>student</i>	<i>course</i>	<i>advisor</i>
<u><i>dept_name</i></u> <i>building</i> <i>budget</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>salary</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u>	<u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>building</i> <i>room_number</i> <i>time_slot_id</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>grade</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>tot_cred</i>	<u><i>course_id</i></u> <i>title</i> <i>dept_name</i> <i>credits</i>	<i>s_id</i> <i>i_id</i>

Definition of “some” Clause

- $F <\text{comp}> \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F <\text{comp}> t)$

Where $<\text{comp}>$ can be: $<$, \leq , $>$, $=$, \neq

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{true}$ (read: 5 < some tuple in the relation)

$(5 < \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{false}$

$(5 = \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true}$

$(5 \neq \text{some } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline \end{array}) = \text{true (since } 0 \neq 5)$

$(= \text{some}) \equiv \text{in}$

However, $(\neq \text{some}) \not\equiv \text{not in}$

Definition of “all” Clause

- $F \text{ <comp> all } r \Leftrightarrow \forall t \in r (F \text{ <comp> } t)$

$$(5 < \text{all } \begin{array}{|c|} \hline 0 \\ \hline 5 \\ \hline 6 \\ \hline \end{array}) = \text{false}$$

$$(5 < \text{all } \begin{array}{|c|} \hline 6 \\ \hline 10 \\ \hline \end{array}) = \text{true}$$

$$(5 = \text{all } \begin{array}{|c|} \hline 4 \\ \hline 5 \\ \hline \end{array}) = \text{false}$$

$$(5 \neq \text{all } \begin{array}{|c|} \hline 4 \\ \hline 6 \\ \hline \end{array}) = \text{true (since } 5 \neq 4 \text{ and } 5 \neq 6)$$

$(\neq \text{all}) \equiv \text{not in}$

However, $(= \text{all}) \not\equiv \text{in}$

Subquery in Where Clause

- The **exists** construct returns the value **true** if the argument subquery is nonempty.
 - Find all courses taught in both the Fall 2017 semester and in the Spring 2018 semester
- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- The **unique** construct evaluates to “true” if a given subquery contains no duplicates .
 - Find all courses that were offered at most once in 2017

				section				
department	instructor	teaches			takes	student	course	
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	
building	name	<u>sec_id</u>	<u>sec_id</u>	<u>sec_id</u>	<u>course_id</u>	name	title	advisor
budget	dept_name	<u>semester</u>	<u>semester</u>	<u>year</u>	<u>sec_id</u>	dept_name	dept_name	<u>s_id</u>
	salary	<u>year</u>	<u>year</u>	building	<u>semester</u>	tot_cred	credits	<u>i_id</u>
				room_number	<u>year</u>			
				time_slot_id	grade			

Query Quest

1. Find names of instructors with salary greater than that of some (at least one) instructor in the Computer Science department.
 - Use self-join, some, exists, aggregation
2. Find all instructors earning the highest salary (there may be more than one with the same salary).
 - Use self-join, all, exists, aggregation

				section				
department	instructor	teaches			takes	student	course	advisor
<u>dept_name</u> building budget	<u>ID</u> name dept_name salary	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>		<u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> building room_number time_slot_id	<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> grade	<u>ID</u> name dept_name tot_cred	<u>course_id</u> title dept_name credits	<u>s_id</u> <u>i_id</u>

Subquery in From and Select Clause

1. Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."
2. Find all departments with the maximum budget (with clause)
3. Find all departments where the total salary is greater than the average of the total salary at all departments
4. List all departments along with the number of instructors in each department

				section				
department	instructor	teaches			takes	student	course	
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>advisor</u>
building	name	<u>sec_id</u>	<u>sec_id</u>	<u>sec_id</u>	<u>course_id</u>	name	title	<u>s_id</u>
budget	dept_name	<u>semester</u>	<u>semester</u>	year	<u>sec_id</u>	dept_name	dept_name	<u>i_id</u>
	salary	<u>year</u>	<u>year</u>	building	<u>semester</u>	tot_cred	credits	
				room_number	<u>year</u>			
				time_slot_id	grade			

With Clause

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs.
- Find all departments with the maximum budget

```
with max_budget (value) as
    (select max(budget)
     from department)
select department.name
from department, max_budget
where department.budget = max_budget.value;
```

<i>department</i>	<i>instructor</i>	<i>teaches</i>	<i>section</i>	<i>takes</i>	<i>student</i>	<i>course</i>	<i>advisor</i>
<u><i>dept_name</i></u> <i>building</i> <i>budget</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>salary</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u>	<u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>building</i> <i>room_number</i> <i>time_slot_id</i>	<u><i>ID</i></u> <u><i>course_id</i></u> <u><i>sec_id</i></u> <u><i>semester</i></u> <u><i>year</i></u> <i>grade</i>	<u><i>ID</i></u> <i>name</i> <i>dept_name</i> <i>tot_cred</i>	<u><i>course_id</i></u> <i>title</i> <i>dept_name</i> <i>credits</i>	<u><i>s_id</i></u> <u><i>i_id</i></u>

Query Quest

- Find all departments where the total salary is greater than the average of the total salary at all departments

department		instructor		teaches		section		takes		student		course		advisor	
<u>dept_name</u>		<u>ID</u>		<u>ID</u>		<u>course_id</u>		<u>ID</u>		<u>ID</u>		<u>course_id</u>		<u>s_id</u>	
building		name		course_id		sec_id		course_id		name		title		i_id	
budget		dept_name		sec_id		semester		sec_id		dept_name		dept_name			
		salary		semester		building		semester		tot_cred		credits			
				year		room_number		year							
						time_slot_id		grade							

Modification of the Database

- Deletion of tuples from a given relation.
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation

```
delete from instructor  
where dept name in (select dept name  
                        from department  
                        where building = 'Watson');
```

department		instructor		teaches		section		takes		student		course		advisor	
<u>dept_name</u> building budget		<u>ID</u> name dept_name salary		<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u>		<u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> building room_number time_slot_id		<u>ID</u> <u>course_id</u> <u>sec_id</u> <u>semester</u> <u>year</u> grade		<u>ID</u> name dept_name tot_cred		<u>course_id</u> title dept_name credits		<u>s_id</u> <u>i_id</u>	

Deletion (Cont.)

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor
where salary < (select avg (salary)
                  from instructor);
```

- Problem: as we delete tuples from *instructor*, the average salary changes
- Solution used in SQL:
 1. First, compute **avg** (salary) and find all tuples to delete
 2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

department		instructor	teaches	section	takes	student	course	advisor
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	<u>s_id</u>
building	name	<u>sec_id</u>	<u>sec_id</u>	<u>semester</u>	<u>course_id</u>	name	title	<u>i_id</u>
budget	dept_name	<u>semester</u>	<u>year</u>	<u>year</u>	<u>sec_id</u>	dept_name	dept_name	
	salary	<u>year</u>	<u>building</u>	<u>room_number</u>	<u>semester</u>	tot_cred	credits	
			<u>time_slot_id</u>	<u>year</u>	<u>grade</u>			

Insertion

- Add a new tuple to *course*

```
insert into course
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently

```
insert into course (course_id, title, dept_name, credits)
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to *student* with *tot_creds* set to null

```
insert into student
values ('3003', 'Green', 'Finance', null);
```

department		instructor		teaches		section		takes		student		course		advisor	
<u>dept_name</u>		<u>ID</u>		<u>course_id</u>		<u>course_id</u>		<u>ID</u>		<u>ID</u>		<u>course_id</u>		<u>s_id</u>	
building		name		sec_id		sec_id		course_id		name		title		i_id	
budget		dept_name		semester		semester		sec_id		dept_name		dept_name			
		salary		year		year		year		tot_cred		credits			
						building		grade							
						room_number									
						time_slot_id									

Insertion (Cont.)

- Make each student in the Music department who has earned more than 144 credit hours an instructor in the Music department with a salary of \$18,000.

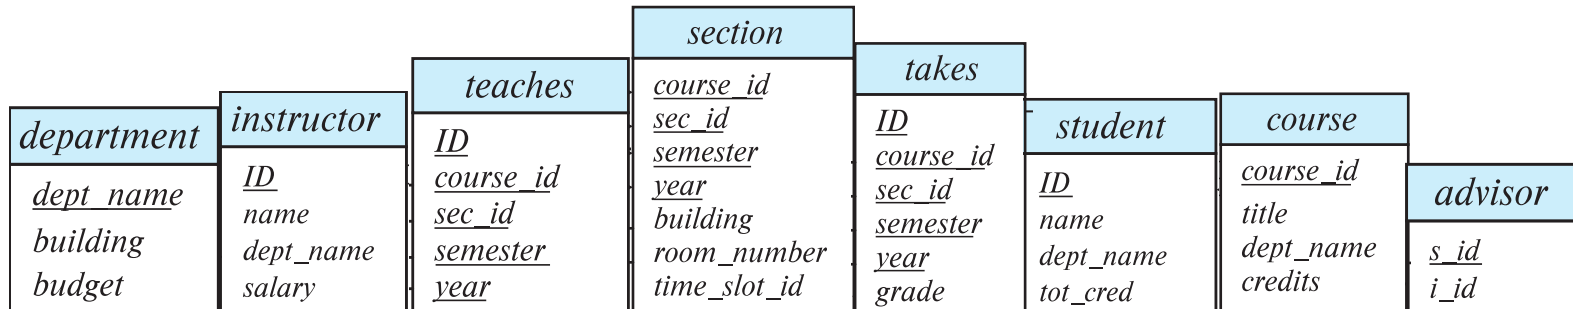
```
insert into instructor
  select ID, name, dept_name, 18000
 from student
 where dept_name = 'Music' and total_cred > 144;
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

Otherwise queries like

```
insert into table1 select * from table1
```

would cause problem



Updates

- Give a 5% salary raise to all instructors

```
update instructor
  set salary = salary * 1.05
```

- Give a 5% salary raise to those instructors who earn less than 70000

```
update instructor
  set salary = salary * 1.05
 where salary < 70000;
```

- Give a 5% salary raise to instructors whose salary is less than average

```
update instructor
  set salary = salary * 1.05
 where salary < (select avg (salary)
                  from instructor);
```

department		instructor		teaches	section	takes	student	course	advisor
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>name</u>	<u>ID</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	
building	dept_name	name	dept_name	<u>course_id</u>	<u>sec_id</u>	<u>course_id</u>	<u>sec_id</u>	title	<u>s_id</u>
budget	salary	dept_name	salary	<u>semester</u>	<u>year</u>	<u>semester</u>	<u>year</u>	dept_name	<u>i_id</u>
				<u>building</u>	<u>room_number</u>	<u>grade</u>	<u>tot_cred</u>	credits	
				<u>time_slot_id</u>					

Updates (Cont.)

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
 - Write two **update** statements:


```
update instructor
  set salary = salary * 1.03
  where salary > 100000;
update instructor
  set salary = salary * 1.05
  where salary <= 100000;
```
 - The order is important
 - Can be done better using the **case** statement (next slide)

department		instructor		teaches	section	takes	student	course	advisor
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>name</u>	<u>ID</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	
building	dept_name	name	dept_name	<u>course_id</u>	<u>sec_id</u>	<u>course_id</u>	<u>sec_id</u>	title	<u>s_id</u>
budget	salary	semester	year	<u>semester</u>	<u>year</u>	<u>semester</u>	<u>year</u>	dept_name	<u>i_id</u>
				<u>year</u>	building	<u>year</u>	tot_cred	credits	
					room_number	grade			
					time_slot_id				

Case Statement for Conditional Updates

- Same query as before but with case statement

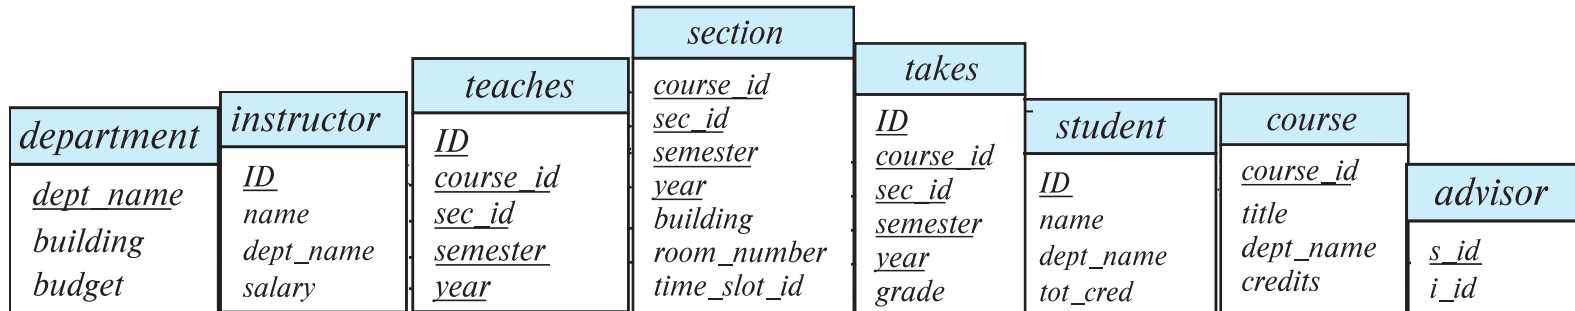
update instructor

set salary = case

when salary <= 100000 then salary * 1.05

else salary * 1.03

end



Updates with Scalar Subqueries

- Recompute and update `tot_creds` value for all students

```

update student S
set tot_cred = (select sum(credits)
                from takes, course
                where takes.course_id = course.course_id and
                    S.ID= takes.ID.and
                    takes.grade <> 'F' and
                    takes.grade is not null);
    
```

- Sets `tot_creds` to null for students who have not taken any course
- Instead of `sum(credits)`, use:

```

case
    when sum(credits) is not null then sum(credits)
    else 0
end
    
```

department		instructor		teaches	section	takes	student	course	advisor
<u>dept_name</u>	<u>ID</u>	<u>ID</u>	<u>name</u>	<u>ID</u>	<u>course_id</u>	<u>ID</u>	<u>ID</u>	<u>course_id</u>	
building	dept_name	name	dept_name	course_id	sec_id	course_id	name	title	s_id
budget	salary	sec_id		semester	year	sec_id	dept_name	dept_name	i_id
		year		building	room_number	semester	tot_cred	credits	
				time_slot_id		year			
						grade			

FIN

Any questions?